



## IBM Developer SKILLS NETWORK

### Introduction Notebook

Estimated time needed: **10** minutes

### Objectives

After completing this lab you will be able to:

- Acquire data in various ways
- Obtain insights from data with Pandas library

### Table of Contents

1. [Data Acquisition \(https:#data\\_acquisition\)](#)
2. [Basic Insight of Dataset \(https:#basic\\_insight\)](#)

## Data Acquisition

There are various formats for a dataset: .csv, .json, .xlsx etc. The dataset can be stored in different places, on your local machine or sometimes online.

In this section, you will learn how to load a dataset into our Jupyter Notebook.

In our case, the Automobile Dataset is an online source, and it is in a CSV (comma separated value) format. Let's use this dataset as an example to practice data reading.

- Data source: [https://archive.ics.uci.edu/ml/machine-learning-databases/autos/imports-85.data?utm\\_medium=Exinfluencer&utm\\_source=Exinfluencer&utm\\_content=000026UJ&utm\\_term=100SkillsNetwork-Channel-SkillsNetworkCoursesIBMDDeveloperSkillsNetworkDA0101ENSkillsNetwork20235326-2022-01-01](https://archive.ics.uci.edu/ml/machine-learning-databases/autos/imports-85.data?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=100SkillsNetwork-Channel-SkillsNetworkCoursesIBMDDeveloperSkillsNetworkDA0101ENSkillsNetwork20235326-2022-01-01)
- Data type: csv

The Pandas Library is a useful tool that enables us to read various datasets into a dataframe; our Jupyter notebook platforms have a built-in **Pandas Library** so that all we need to do is import Pandas without installing.

```
In [ ]: #install specific version of libraries used in lab
        #! mamba install pandas==1.3.3 -y
        #! mamba install numpy=1.21.2 -y
```

```
In [1]: # import pandas library
        import pandas as pd
        import numpy as np
```

## Read Data

We use `pandas.read_csv()` function to read the csv file. In the brackets, we put the file path along with a quotation mark so that pandas will read the file into a dataframe from that address. The file path can be either an URL or your local file address.

Because the data does not include headers, we can add an argument `headers = None` inside the `read_csv()` method so that pandas will not automatically set the first row as a header.

You can also assign the dataset to any variable you create.

This dataset was hosted on IBM Cloud object. Click [HERE](https://cocl.us/DA101EN_object_storage?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=100065SkillsNetwork-Channel-SkillsNetworkCoursesIBMDDeveloperSkillsNetworkDA0101ENSkillsNetwork20235326-2022-01-01) ([https://cocl.us/DA101EN\\_object\\_storage?utm\\_medium=Exinfluencer&utm\\_source=Exinfluencer&utm\\_content=000026UJ&utm\\_term=100065SkillsNetwork-Channel-SkillsNetworkCoursesIBMDDeveloperSkillsNetworkDA0101ENSkillsNetwork20235326-2022-01-01](https://cocl.us/DA101EN_object_storage?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=100065SkillsNetwork-Channel-SkillsNetworkCoursesIBMDDeveloperSkillsNetworkDA0101ENSkillsNetwork20235326-2022-01-01)) for free storage.

```
In [2]: # Import pandas library
        import pandas as pd

        # Read the online file by the URL provides above, and assign it to variable "df"
        other_path = "https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/"
        df = pd.read_csv(other_path, header=None)
```

After reading the dataset, we can use the `dataframe.head(n)` method to check the top `n` rows of the dataframe, where `n` is an integer. Contrary to `dataframe.head(n)`, `dataframe.tail(n)` will show you the bottom `n` rows of the dataframe.

```
In [3]: # show the first 5 rows using dataframe.head() method
print("The first 5 rows of the dataframe")
df.head(5)
```

The first 5 rows of the dataframe

Out[3]:

	0	1	2	3	4	5	6	7	8	9	...	16	17	18	19	20	21
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0	111
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0	111
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...	152	mpfi	2.68	3.47	9.0	154
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	...	109	mpfi	3.19	3.40	10.0	102
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	...	136	mpfi	3.19	3.40	8.0	115

5 rows × 26 columns



`df.tail(10)`

## Question #1:

Check the bottom 10 rows of data frame "df".

In [4]: *# Write your code below and press Shift+Enter to execute*  
`df.tail(10)`

Out[4]:

	0	1	2	3	4	5	6	7	8	9	...	16	17	18	19	20	
<b>195</b>	-1	74	volvo	gas	std	four	wagon	rwd	front	104.3	...	141	mpfi	3.78	3.15	9.5	1
<b>196</b>	-2	103	volvo	gas	std	four	sedan	rwd	front	104.3	...	141	mpfi	3.78	3.15	9.5	1
<b>197</b>	-1	74	volvo	gas	std	four	wagon	rwd	front	104.3	...	141	mpfi	3.78	3.15	9.5	1
<b>198</b>	-2	103	volvo	gas	turbo	four	sedan	rwd	front	104.3	...	130	mpfi	3.62	3.15	7.5	1
<b>199</b>	-1	74	volvo	gas	turbo	four	wagon	rwd	front	104.3	...	130	mpfi	3.62	3.15	7.5	1
<b>200</b>	-1	95	volvo	gas	std	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	9.5	1
<b>201</b>	-1	95	volvo	gas	turbo	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	8.7	1
<b>202</b>	-1	95	volvo	gas	std	four	sedan	rwd	front	109.1	...	173	mpfi	3.58	2.87	8.8	1
<b>203</b>	-1	95	volvo	diesel	turbo	four	sedan	rwd	front	109.1	...	145	idi	3.01	3.40	23.0	1
<b>204</b>	-1	95	volvo	gas	turbo	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	9.5	1

10 rows × 26 columns



[Click here for the solution](#)

## Add Headers

Take a look at our dataset. Pandas automatically set the header with an integer starting from 0.

To better describe our data, we can introduce a header. This information is available at:

<https://archive.ics.uci.edu/ml/datasets/Automobile>

([https://archive.ics.uci.edu/ml/datasets/Automobile?](https://archive.ics.uci.edu/ml/datasets/Automobile?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=1000655)

[utm\\_medium=Exinfluencer&utm\\_source=Exinfluencer&utm\\_content=000026UJ&utm\\_term=1000655](https://archive.ics.uci.edu/ml/datasets/Automobile?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=1000655)

[SkillsNetwork-Channel-](https://archive.ics.uci.edu/ml/datasets/Automobile?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=1000655)

[SkillsNetworkCoursesIBMDDeveloperSkillsNetworkDA0101ENSkillsNetwork20235326-2022-01-01\).](https://archive.ics.uci.edu/ml/datasets/Automobile?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=1000655)

Thus, we have to add headers manually.

First, we create a list "headers" that include all column names in order. Then, we use `dataframe.columns = headers` to replace the headers with the list we created.



```
In [5]: # create headers list
headers = ["symboling", "normalized-losses", "make", "fuel-type", "aspiration", "num-
          "drive-wheels", "engine-location", "wheel-base", "length", "width", "height",
          "num-of-cylinders", "engine-size", "fuel-system", "bore", "stroke", "compression-
          "peak-rpm", "city-mpg", "highway-mpg", "price"]
print("headers\n", headers)
```

```
headers
['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration', 'num-of-
doors', 'body-style', 'drive-wheels', 'engine-location', 'wheel-base', 'length',
'width', 'height', 'curb-weight', 'engine-type', 'num-of-cylinders', 'engine-size',
'fuel-system', 'bore', 'stroke', 'compression-ratio', 'horsepower', 'peak-rpm',
'city-mpg', 'highway-mpg', 'price']
```

We replace headers and recheck our dataframe:

```
In [6]: df.columns = headers
df.head(10)
```

Out[6]:

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front	94.5
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4
5	2	?	audi	gas	std	two	sedan	fwd	front	99.8
6	1	158	audi	gas	std	four	sedan	fwd	front	105.8
7	1	?	audi	gas	std	four	wagon	fwd	front	105.8
8	1	158	audi	gas	turbo	four	sedan	fwd	front	105.8
9	0	?	audi	gas	turbo	two	hatchback	4wd	front	99.5

10 rows × 26 columns



We need to replace the "?" symbol with NaN so the dropna() can remove the missing values:

```
In [9]: df1=df.replace('?',np.NaN)
df1
```

Out[9]:

	symboling	normalized- losses	make	fuel- type	aspiration	num- of- doors	body- style	drive- wheels	engine- location	wheel base
0	3	NaN	alfa-romero	gas	std	two	convertible	rwd	front	88.6
1	3	NaN	alfa-romero	gas	std	two	convertible	rwd	front	88.6
2	1	NaN	alfa-romero	gas	std	two	hatchback	rwd	front	94.4
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4
...	...	...	...	...	...	...	...	...	...	...
200	-1	95	volvo	gas	std	four	sedan	rwd	front	109.1
201	-1	95	volvo	gas	turbo	four	sedan	rwd	front	109.1
202	-1	95	volvo	gas	std	four	sedan	rwd	front	109.1
203	-1	95	volvo	diesel	turbo	four	sedan	rwd	front	109.1
204	-1	95	volvo	gas	turbo	four	sedan	rwd	front	109.1

201 rows × 26 columns



We can drop missing values along the column "price" as follows:

```
In [8]: df=df1.dropna(subset=["price"], axis=0)
df.head(20)
```

Out[8]:

	symboling	normalized- losses	make	fuel- type	aspiration	num- of- doors	body- style	drive- wheels	engine- location	wheel- base
0	3	NaN	alfa-romero	gas	std	two	convertible	rwd	front	88.6
1	3	NaN	alfa-romero	gas	std	two	convertible	rwd	front	88.6
2	1	NaN	alfa-romero	gas	std	two	hatchback	rwd	front	94.5
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4
5	2	NaN	audi	gas	std	two	sedan	fwd	front	99.8
6	1	158	audi	gas	std	four	sedan	fwd	front	105.8
7	1	NaN	audi	gas	std	four	wagon	fwd	front	105.8
8	1	158	audi	gas	turbo	four	sedan	fwd	front	105.8
10	2	192	bmw	gas	std	two	sedan	rwd	front	101.2
11	0	192	bmw	gas	std	four	sedan	rwd	front	101.2
12	0	188	bmw	gas	std	two	sedan	rwd	front	101.2
13	0	188	bmw	gas	std	four	sedan	rwd	front	101.2
14	1	NaN	bmw	gas	std	four	sedan	rwd	front	103.5
15	0	NaN	bmw	gas	std	four	sedan	rwd	front	103.5
16	0	NaN	bmw	gas	std	two	sedan	rwd	front	103.5
17	0	NaN	bmw	gas	std	four	sedan	rwd	front	110.0
18	2	121	chevrolet	gas	std	two	hatchback	fwd	front	88.4
19	1	98	chevrolet	gas	std	two	hatchback	fwd	front	94.5
20	0	81	chevrolet	gas	std	four	sedan	fwd	front	94.5

20 rows × 26 columns



Now, we have successfully read the raw dataset and added the correct headers into the dataframe.

## Question #2:

Find the name of the columns of the dataframe.

In [10]: *# Write your code below and press Shift+Enter to execute*  
df.columns

Out[10]: Index(['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration',  
'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',  
'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',  
'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',  
'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',  
'highway-mpg', 'price'],  
dtype='object')

[Click here for the solution](#)

## Save Dataset

Correspondingly, Pandas enables us to save the dataset to csv. By using the `dataframe.to_csv()` method, you can add the file path and name along with quotation marks in the brackets.

For example, if you would save the dataframe **df** as **automobile.csv** to your local machine, you may use the syntax below, where `index = False` means the row names will not be written.

```
df.to_csv("automobile.csv", index=False)
```

We can also read and save other file formats. We can use similar functions like `pd.read_csv()` and `df.to_csv()` for other data formats. The functions are listed in the following table:

## Read/Save Other Data Formats

Data Formate	Read	Save
csv	<code>pd.read_csv()</code>	<code>df.to_csv()</code>
json	<code>pd.read_json()</code>	<code>df.to_json()</code>
excel	<code>pd.read_excel()</code>	<code>df.to_excel()</code>
hdf	<code>pd.read_hdf()</code>	<code>df.to_hdf()</code>
sql	<code>pd.read_sql()</code>	<code>df.to_sql()</code>
...	...	...

## Basic Insight of Dataset

After reading data into Pandas dataframe, it is time for us to explore the dataset.



There are several ways to obtain essential insights of the data to help us better understand our dataset.

## Data Types

Data has a variety of types.

The main types stored in Pandas dataframes are **object**, **float**, **int**, **bool** and **datetime64**. In order to better learn about each attribute, it is always good for us to know the data type of each column. In Pandas:

```
In [11]: df.dtypes
```

```
Out[11]: symboling          int64
normalized-losses    object
make                 object
fuel-type            object
aspiration           object
num-of-doors         object
body-style           object
drive-wheels         object
engine-location      object
wheel-base          float64
length              float64
width               float64
height              float64
curb-weight          int64
engine-type          object
num-of-cylinders     object
engine-size          int64
fuel-system          object
bore                 object
stroke              object
compression-ratio    float64
horsepower           object
peak-rpm             object
city-mpg             int64
highway-mpg          int64
price                object
dtype: object
```

A series with the data type of each column is returned.

```
In [14]: # check the data type of data frame "df" by .dtypes
print(df.dtypes)
```

```
symboling          int64
normalized-losses  object
make              object
fuel-type          object
aspiration         object
num-of-doors       object
body-style         object
drive-wheels       object
engine-location    object
wheel-base        float64
length            float64
width             float64
height            float64
curb-weight        int64
engine-type        object
num-of-cylinders   object
engine-size        int64
fuel-system        object
bore              object
stroke            object
compression-ratio  float64
horsepower         object
peak-rpm          object
city-mpg           int64
highway-mpg        int64
price             object
dtype: object
```

As shown above, it is clear to see that the data type of "symboling" and "curb-weight" are `int64` , "normalized-losses" is `object` , and "wheel-base" is `float64` , etc.

These data types can be changed; we will learn how to accomplish this in a later module.

## Describe

If we would like to get a statistical summary of each column e.g. count, column mean value, column standard deviation, etc., we use the describe method:

```
dataframe.describe()
```

This method will provide various summary statistics, excluding `NaN` (Not a Number) values.

In [15]: `df.describe()`

Out[15]:

	symboling	wheel- base	length	width	height	curb-weight	engine- size	comp
<b>count</b>	201.000000	201.000000	201.000000	201.000000	201.000000	201.000000	201.000000	20
<b>mean</b>	0.840796	98.797015	174.200995	65.889055	53.766667	2555.666667	126.875622	1
<b>std</b>	1.254802	6.066366	12.322175	2.101471	2.447822	517.296727	41.546834	
<b>min</b>	-2.000000	86.600000	141.100000	60.300000	47.800000	1488.000000	61.000000	
<b>25%</b>	0.000000	94.500000	166.800000	64.100000	52.000000	2169.000000	98.000000	
<b>50%</b>	1.000000	97.000000	173.200000	65.500000	54.100000	2414.000000	120.000000	
<b>75%</b>	2.000000	102.400000	183.500000	66.600000	55.500000	2926.000000	141.000000	
<b>max</b>	3.000000	120.900000	208.100000	72.000000	59.800000	4066.000000	326.000000	2

This shows the statistical summary of all numeric-typed (int, float) columns.

For example, the attribute "symboling" has 205 counts, the mean value of this column is 0.83, the standard deviation is 1.25, the minimum value is -2, 25th percentile is 0, 50th percentile is 1, 75th percentile is 2, and the maximum value is 3.

However, what if we would also like to check all the columns including those that are of type object?


You can add an argument `include = "all"` inside the bracket. Let's try it again.

```
In [16]: # describe all the columns in "df"
df.describe(include = "all")
```

Out[16]:

	symboling	normalized- losses	make	fuel- type	aspiration	num- of- doors	body- style	drive- wheels	engine- location	whe ba
<b>count</b>	201.000000	164	201	201	201	199	201	201	201	201.0000
<b>unique</b>	NaN	51	22	2	2	2	5	3	2	NaN
<b>top</b>	NaN	161	toyota	gas	std	four	sedan	fwd	front	NaN
<b>freq</b>	NaN	11	32	181	165	113	94	118	198	NaN
<b>mean</b>	0.840796	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	98.7970
<b>std</b>	1.254802	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	6.0663
<b>min</b>	-2.000000	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	86.6000
<b>25%</b>	0.000000	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	94.5000
<b>50%</b>	1.000000	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	97.0000
<b>75%</b>	2.000000	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	102.4000
<b>max</b>	3.000000	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	120.9000

11 rows × 26 columns



Now it provides the statistical summary of all the columns, including object-typed attributes.

We can now see how many unique values there, which one is the top value and the frequency of top value in the object-typed columns.

Some values in the table above show as "NaN". This is because those numbers are not available regarding a particular column type.

## Question #3:

You can select the columns of a dataframe by indicating the name of each column. For example, you can select the three columns as follows:

```
dataframe[['column 1 ', 'column 2', 'column 3']]
```

Where "column" is the name of the column, you can apply the method ".describe()" to get the statistics of those columns as follows:

```
dataframe[['column 1 ', 'column 2', 'column 3']].describe()
```

Apply the method to ".describe()" to the columns 'length' and 'compression-ratio'.

```
In [17]: # Write your code below and press Shift+Enter to execute
df[['make','price']]
```

Out[17]:

	make	price
0	alfa-romero	13495
1	alfa-romero	16500
2	alfa-romero	16500
3	audi	13950
4	audi	17450
...	...	...
200	volvo	16845
201	volvo	19045
202	volvo	21485
203	volvo	22470
204	volvo	22625

201 rows × 2 columns

[Click here for the solution](#)

## Info

Another method you can use to check your dataset is:

```
dataframe.info()
```

It provides a concise summary of your DataFrame.

This method prints information about a DataFrame including the index dtype and columns, non-null values and memory usage.

```
In [18]: # Look at the info of "df"
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 201 entries, 0 to 204
Data columns (total 26 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   symboling           201 non-null    int64
 1   normalized-losses   164 non-null    object
 2   make                201 non-null    object
 3   fuel-type           201 non-null    object
 4   aspiration           201 non-null    object
 5   num-of-doors        199 non-null    object
 6   body-style          201 non-null    object
 7   drive-wheels        201 non-null    object
 8   engine-location     201 non-null    object
 9   wheel-base         201 non-null    float64
10   length              201 non-null    float64
11   width               201 non-null    float64
12   height              201 non-null    float64
13   curb-weight         201 non-null    int64
14   engine-type         201 non-null    object
15   num-of-cylinders    201 non-null    object
16   engine-size         201 non-null    int64
17   fuel-system         201 non-null    object
18   bore                197 non-null    object
19   stroke              197 non-null    object
20   compression-ratio   201 non-null    float64
21   horsepower          199 non-null    object
22   peak-rpm            199 non-null    object
23   city-mpg            201 non-null    int64
24   highway-mpg         201 non-null    int64
25   price               201 non-null    object
dtypes: float64(5), int64(5), object(16)
memory usage: 42.4+ KB
```

## Excellent! You have just completed the Introduction Notebook!

Thank you for completing this lab!

### Author

[Joseph Santarcangelo \(https://www.linkedin.com/in/joseph-s-50398b136/?utm\\_medium=Exinfluencer&utm\\_source=Exinfluencer&utm\\_content=000026UJ&utm\\_term=100065f](https://www.linkedin.com/in/joseph-s-50398b136/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=100065f)  
[SkillsNetwork-Channel-](#)  
[SkillsNetworkCoursesIBMDDeveloperSkillsNetworkDA0101ENSkillsNetwork20235326-2022-01-01\)](#)

### Other Contributors

[Mahdi Noorian PhD \(https://www.linkedin.com/in/mahdi-noorian-58219234/?utm\\_medium=Exinfluencer&utm\\_source=Exinfluencer&utm\\_content=000026UJ&utm\\_term=1000655SkillsNetwork-Channel-SkillsNetworkCoursesIBMDDeveloperSkillsNetworkDA0101ENSkillsNetwork20235326-2022-01-01\)](https://www.linkedin.com/in/mahdi-noorian-58219234/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=1000655SkillsNetwork-Channel-SkillsNetworkCoursesIBMDDeveloperSkillsNetworkDA0101ENSkillsNetwork20235326-2022-01-01)

Bahare Talayian

Eric Xiao

Steven Dong

Parizad

Hima Vasudevan

[Fiorella Wenver \(https://www.linkedin.com/in/fiorellawever/?utm\\_medium=Exinfluencer&utm\\_source=Exinfluencer&utm\\_content=000026UJ&utm\\_term=1000655SkillsNetwork-Channel-SkillsNetworkCoursesIBMDDeveloperSkillsNetworkDA0101ENSkillsNetwork20235326-2022-01-01\)](https://www.linkedin.com/in/fiorellawever/?utm_medium=Exinfluencer&utm_source=Exinfluencer&utm_content=000026UJ&utm_term=1000655SkillsNetwork-Channel-SkillsNetworkCoursesIBMDDeveloperSkillsNetworkDA0101ENSkillsNetwork20235326-2022-01-01)

[Yi Yao \(https://www.linkedin.com/in/yi-leng-yao-84451275/\)](https://www.linkedin.com/in/yi-leng-yao-84451275/)

## Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2020-10-30	2.3	Lakshmi	Changed URL of the csv
2020-09-22	2.2	Nayef	Added replace() method to remove '?'
2020-09-09	2.1	Lakshmi	Made changes in info method of dataframe
2020-08-27	2.0	Lavanya	Moved lab to course repo in GitLab

© IBM Corporation 2020. All rights reserved.

